SECOND ORDER

AI, Capital & Work · A Razor Research client briefing

ISSUE 02

June 10, 2026

THE WRONG DENOMINATOR

The market just discovered it has been measuring AI in the wrong units. What the shift from tokens to cost per completed task means for budgets, valuations, and your operating model.

THE CLAIM

Token consumption is the metric that doubled the semiconductor complex and underpins two near-trillion-dollar IPO runways. It measures revenue to the labs, not value to the buyer. Measured per completed task, the cost of intelligence is collapsing. Both facts are true at once, and the gap between them is where the next 18 months of corporate strategy and market repricing will play out.

- Apollo co-president John Zito, June 10: tokenmaxxing is *"a lot of BS, honestly… prices are collapsing per unit of IQ."* Goldman's desks landed on the same frame within the week: useful task completion per watt and per dollar.
- The bills are exploding anyway. Average business token spend is up about 13x since January 2025 (Ramp), and a typical agent job burns about 96,000 tokens (SemiAnalysis). Flat-rate pricing ended June 1, when GitHub moved every Copilot plan to usage billing. Anthropic filed its confidential S-1 the same morning.
- Nobody is cutting AI. In UBS's enterprise checks, about 60% call token costs a real issue, and effectively none are retreating. They cut **around** AI instead: other IT spend, cloud, and headcount growth (Exhibit 4).
- The second-order trade: if good enough will do, volume migrates to cheap and open models, smart routing, and edge inference. The dollars pool away from where frontier-priced valuations assume.
- For operators: govern **cost per completed task**, not adoption. Token budgets are starting to work the way headcount budgets always have: capped, forecast, and fought over.

01 · WHAT HAPPENED

Sixty days that repriced AI

For two years the AI trade ran on a simple proxy: more tokens means more intelligence consumed, which means more value created. That proxy broke in the spring. Agents changed the unit economics of usage. An agent does not chat; it loops, plans, retries, and reviews, and a typical job consumes on the order of 96,000 tokens. Providers responded by ending the subsidy. OpenAI moved Codex to usage billing on April 2, Google followed with Gemini on May 19, and on June 1 GitHub converted every Copilot plan to metered "AI credits." Within sixty days the AI cost line went from a rounding error to a board topic (Exhibit 1).

Companies reacted fast, and unevenly. Uber, which had run internal leaderboards encouraging engineers to maximize AI usage, burned its full-year 2026 AI budget in roughly four months and then capped engineers at $1,500 per month per tool. Walmart capped its internal AI agent after demand ran too hot. Microsoft cancelled most of its Claude Code licenses partly on cost. Axios reported one enterprise spending $500 million on Claude in a single month with no usage limits; the figure is single-sourced and unconfirmed, but it fits everything else on the tape. The Wall Street Journal's summary: "the free-money period for AI is definitively over."

EXHIBIT 1

Sixty days from tokenmaxxing to tokenpanic

Key repricing events, April 2 to June 10, 2026

Apr 2	OpenAI Codex moves to usage-based billing
Apr 27	Developer backlash to Copilot repricing: “pay the same, get less”
May 19	Google Gemini moves to usage-based billing
May 26	Uber burns its full-year 2026 AI budget in four months
May 28	Axios “AI sticker shock”; $500M single-month Claude bill reported*
Jun 1	GitHub Copilot: every plan usage-based · Anthropic files confidential S-1
Jun 2	Uber caps tokens at $1,500/engineer/mo; Walmart caps its internal agent
Jun 8	Citrini Research: “tokenmaxxing → tokenpanic”
Jun 10	Apollo’s Zito: price per unit of IQ, not per token: “token talk is BS”

THE PANIC WEEK

*Note: *Single-sourced (Axios, anonymous consultant); unconfirmed by Anthropic.*

Source: Company announcements; Bloomberg; Axios; Fortune; GitHub changelog; Citrini Research.

The tape noticed. The Silicon Data LLM Token Expenditure Index, a usage-weighted measure of what the market actually pays per million tokens, doubled from December to an all-time high on May 20. Then it stalled and printed six consecutive down days, the longest streak since January (Exhibit 2). The reading is ambiguous, and the ambiguity is the story. A falling expenditure index can mean demand is rolling over, or it can mean deflation arriving in the spend line: the same work, bought cheaper. Citadel’s desk leans toward the second. Adoption, they wrote, is becoming “less about what frontier models can do and more about the price.”

EXHIBIT 2

The tape rolled over: token expenditure index stalls after doubling

Silicon Data LLM Token Expenditure Index: usage-weighted market price per M tokens (Dec 2025 = 1.00)

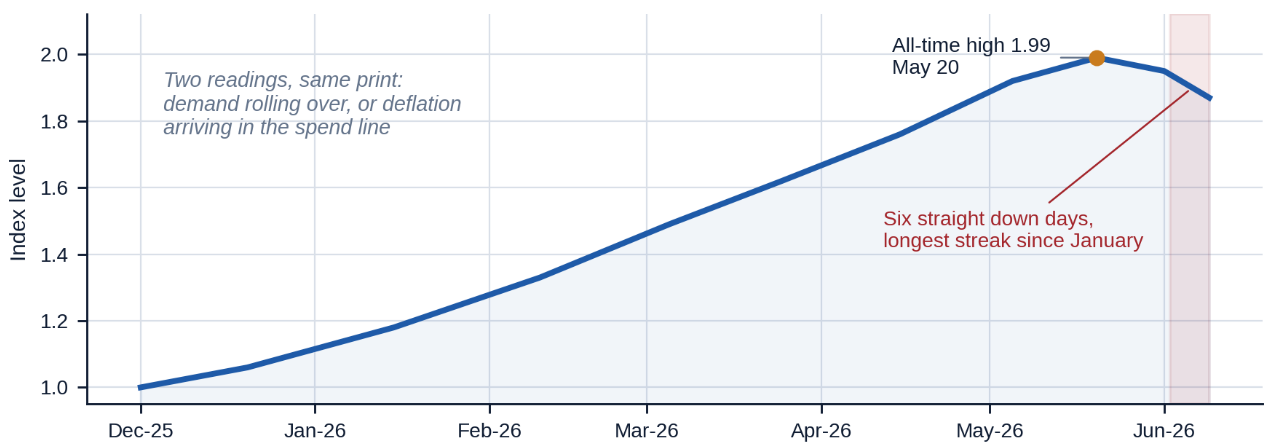

Note: Reported anchors: Dec 2025 base, ~1.99 peak May 20 2026, six consecutive down days through Jun 9. Intra-period path estimated.

Source: Silicon Data Token Market Pulse, reported levels via Seeking Alpha (A. Steno Larsen) and Citadel desk commentary, Jun 2026.

02 · THE REFRAME

The measurement error: tokens price throughput, not intelligence

Into this confusion stepped John Zito, co-president of Apollo, at the Morgan Stanley Financials Conference: *"I think tokenmaxxing and token talk is — it's a lot of BS, honestly. If you look at per unit of knowledge and cost per unit of knowledge, prices are collapsing. Prices are collapsing per unit of IQ, if you did it that way."* The point is more than rhetorical. A token is a unit of throughput, and nobody buys throughput. Price the capability instead, the way a 2026 laptop costs what a 2010 laptop did while doing far more, and the cost of intelligence is in freefall even as the bills explode.

The data supports the denominator. Stanford's AI Index documents a 280-fold collapse in the cost of GPT-3.5-class output in eighteen months, from $20.00 to $0.07 per million tokens. Open-weight and Chinese models now deliver near-frontier capability at a tenth to a twenty-fifth of frontier prices. Cursor's latest model matches frontier coding performance at roughly one tenth the cost per task. Goldman partner Rich Privorotsky framed the correct metric days before Zito: useful task completion per watt and per dollar. When two institutions of that weight land on the same framework inside a week, a consensus is forming (Exhibit 3).

EXHIBIT 3

The scissors: the unit cost of intelligence deflates while the bills inflate

LHS (log): $/M tokens for GPT-3.5-class output · RHS: avg. business token spend, indexed (Jan 2025 = 1)

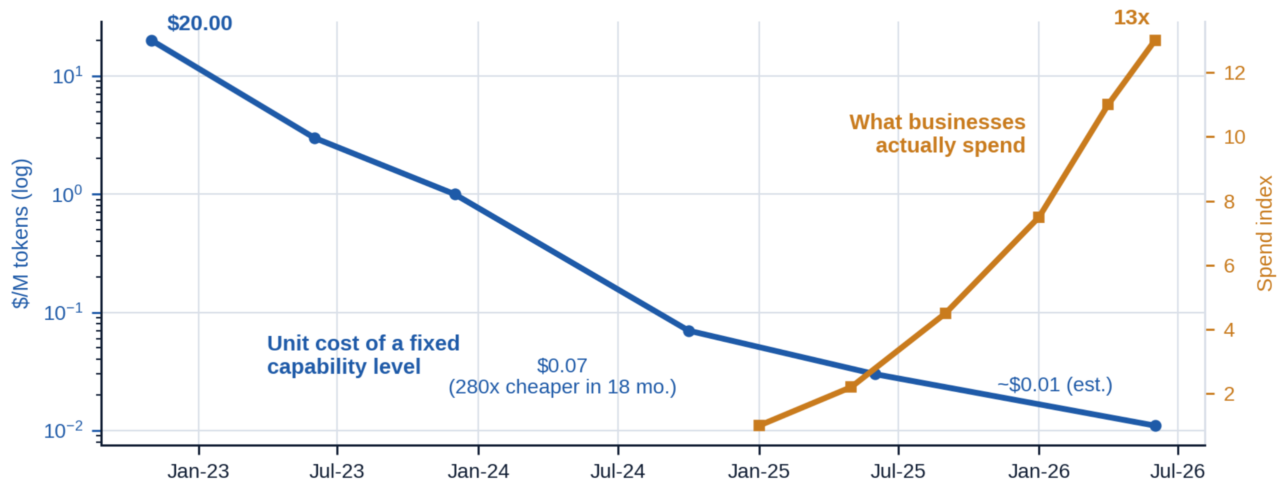

Note: Intermediate points estimated (cost: ~10x/yr deflation trend; spend: smoothed path between reported endpoints).

Source: Stanford AI Index 2025 (cost anchors, Nov 2022 / Oct 2024); Ramp via SemiAnalysis/D. Thompson (spend, Jan 2025 / May 2026).

The CFOs are right about the numerator, though. Gartner finds that even a 90% collapse in inference costs would not make enterprise AI cheaper, because agents consume tokens faster than prices fall and providers do not fully pass savings through. That matches the lived experience of every company in the UBS checks. A collapsing unit cost times an exploding unit count is still a bigger bill. The resolution: token expenditure is a useless productivity metric and a decisive revenue metric. It says nothing about value created. It says a great deal about whether the revenue under roughly $3 trillion of AI-linked valuation is durable. That is why the index rollover matters even if Zito is right about everything else.

03 · THE CORPORATE RESPONSE

Rationing without retreat: the capital swap

The most counterintuitive finding in the UBS Evidence Lab checks is what enterprises are *not* doing. About 60% say token costs have become a real issue ("this is not a made-up media story"). One described the Copilot pricing change in a single word: "chaos." Yet not one check showed a company slamming the brakes. The dominant behavior is guardrails: caps, alerts, model downshifting, pooled token budgets. Several explicitly refused to throttle usage and are cannibalizing other line items to make room, starting with external IT services, cloud spend, and, most notably, headcount growth.

That last item deserves its own exhibit. Gartner-surveyed CFO headcount growth expectations fell from 6% to 2% for 2026 while three quarters of CFOs raise technology budgets, nearly half by double digits. Atlanta Fed survey data shows executives expecting +2.25% productivity from AI over three years against only -1.2% headcount. This is not mass layoffs. It is a quiet, compounding substitution of compute budgets for hiring plans. Walmart has frozen its 2.1 million-person workforce for three years while revenue grows. The planning template spreading across the economy is **flat headcount plus AI**. We call it the capital swap (Exhibit 4).

EXHIBIT 4

The capital swap: compute budgets are eating headcount plans

CFO and executive survey readings, 2025 to 2026

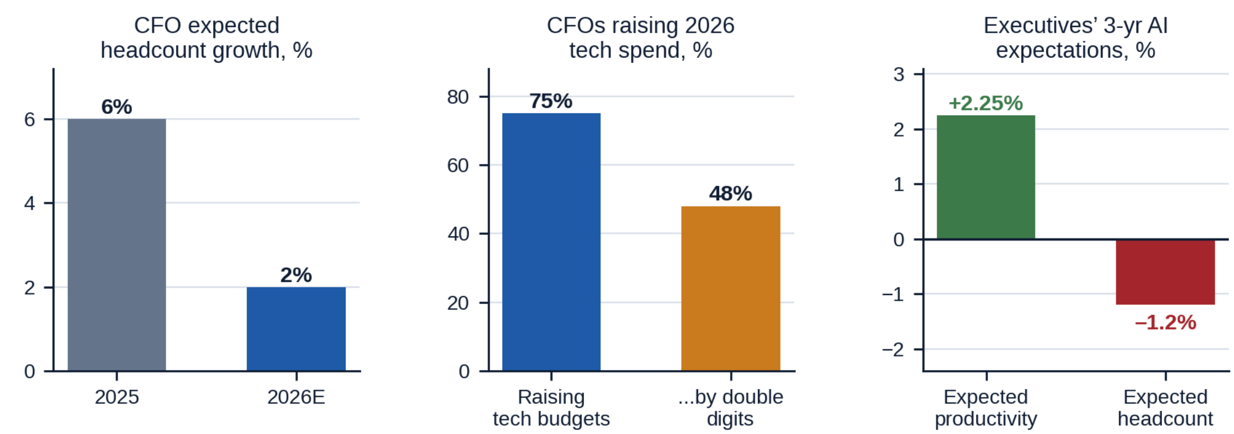

Note: Panels use different surveys and horizons; read as direction, not level.

Source: Gartner CFO surveys via The CFO (Feb 2026); Atlanta Fed Survey of Business Uncertainty macroblog (May 2026).

Structurally, software is turning from a fixed cost into a variable cost. Exponential View's survey work found over 70% of companies blew through their 2025 AI budgets. SemiAnalysis's Doug O'Laughlin argues "automated intelligence" is simply a permanent new opex category, and that firms will swing between overspend and underspend until each finds its own **labor-to-compute ratio**. Boards will start asking for that ratio, not an adoption percentage, by this time next year.

04 · THE MACRO STAKES

A one-trade economy funding a deflationary technology

The denominator debate would be academic if the numbers were small. They are not. Hyperscaler capex roughly doubled in two years and is guided to $600 to $805 billion in 2026, which is about 60% of big-tech operating cash flow and more than double global telecom capex (Exhibit 5). Morgan Stanley estimates AI-related investment accounted for roughly three quarters of Q1 2026 US GDP growth; without it, the economy was approximately flat (Exhibit 6). Apollo's own chief economist, Torsten Slok, calculates the top ten S&P 500 names are more richly valued than at the 1999 peak. The bubble thesis and the deflation thesis now coexist inside the same firm, which tells you the debate is real.

EXHIBIT 5

The denominator problem is funded by the biggest capex cycle in tech history

Hyperscaler capital expenditure, $B (big-5: AMZN, MSFT, GOOGL, META, ORCL)

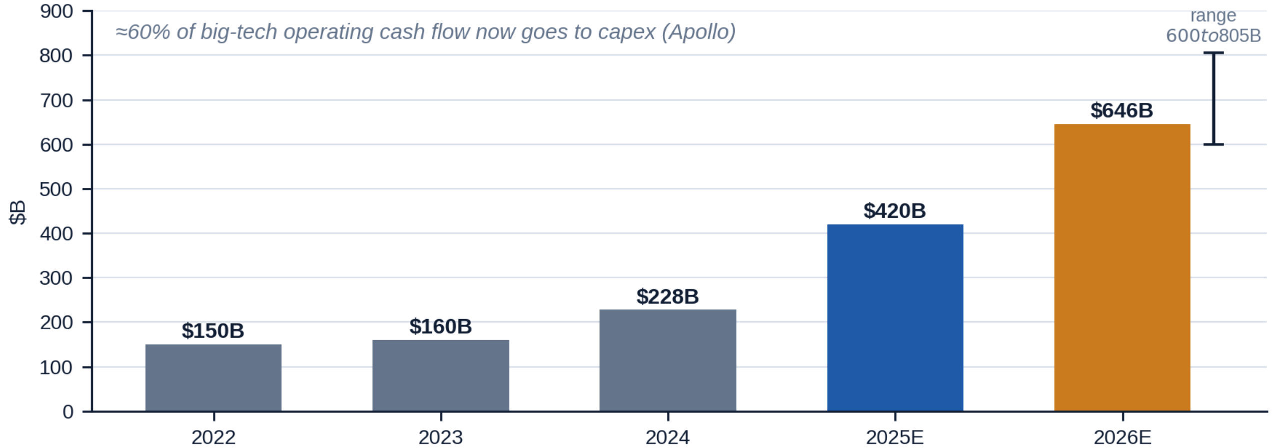

Note: 2022/2025 estimated from compiled filings; definitions vary (big-4 vs big-5). 2026E shown at consensus with range whisker.

Source: Company filings via Apollo/Epoch AI compilations; 2026 range: consensus ~USD 646B (IEEE ComSoc) to USD 805B big-5 (Morgan Stanley).

EXHIBIT 6

Growth is increasingly the buildout itself

Estimated share of US real GDP growth attributable to AI-related investment

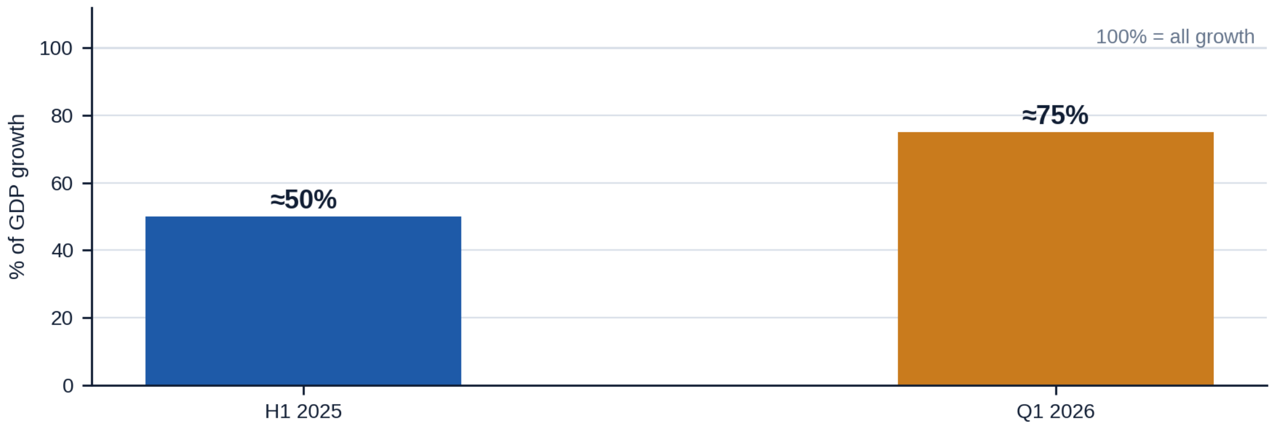

Note: Different estimation methodologies; directionally consistent, not directly comparable. Ex-AI growth ≈ flat in Q1 2026 (MS).

Source: H1 2025: P. Kedrosky; Q1 2026: Morgan Stanley via Fortune (May 2026). St. Louis Fed framework for attribution.

SECOND-ORDER EFFECT: THE DEFLATION PARADOX

Zito's most under-quoted line: *"If AI is real, it's so hyper-deflationary to so many things over the long term that it's really hard to take risk."* The bull case for the technology is a bear case for pricing power: in software, in services, in any business whose margin is rent on cognitive friction.

The timing also inverts intuition. BIS modeling shows that if households and firms *anticipate* AI-driven income gains, AI is inflationary first. Gita Gopinath argues the buildout's appetite for turbines, contractors, and savings is already pushing global rates up. The second-order read: **AI raises the cost of capital before it lowers the cost of anything else.** Leveraged balance sheets should plan accordingly.

05 · SPILLOVER

The narrative is doing the firing

Meanwhile the cost panic is leaking into the labor market through the narrative channel, ahead of the capability channel. Challenger's May report makes “AI” the number one stated reason for US job cuts for the third consecutive month: 38,579 cuts citing AI, 40% of the month's total, up from 7% in January. Year-to-date AI-attributed cuts already exceed all of 2025 (Exhibit 7). This is employer self-attribution, though. HBR's Davenport and Srinivasan document that the wave is **anticipatory**: firms are cutting ahead of demonstrated AI capability because investors reward the story. One CEO put it more sharply: some companies may be citing AI to cover cuts made to *pay for* AI.

EXHIBIT 7

“AI” cited in 40% of May job cuts, up from 7% in January

Share of announced US job cuts attributing AI, %, 2026

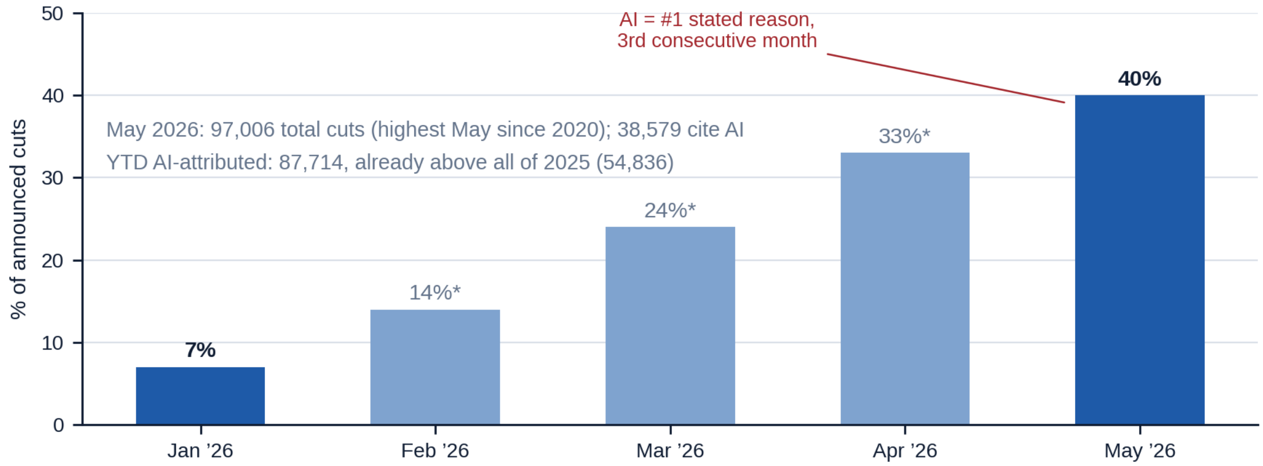

Note: Jan (7%) and May (40%) reported; Feb-Apr (*) interpolated. Self-attribution: “AI” may stand in for ordinary cost cuts.

Source: Challenger, Gray & Christmas May 2026 report (Jun 4, 2026); CNBC.

Recall the floor argument from our first briefing: substitution stops where the marginal cost of compute exceeds the marginal cost of labor for the task. Tokenpanic just made that floor visible. When an agent job costs real money, “automate everything” becomes a budget line someone must defend. The displacement question is genuinely contested in the data; the NY Fed now attributes most of the rise in new-graduate unemployment to remote work, not AI, and that fight is the subject of our next issue. What is not contested: the *story* of AI is restructuring organizations faster than the technology is.

06 · WHERE THE DOLLARS POOL

Good enough will do, and the IPO tension

If the denominator wins, if buyers optimize cost per completed task, the demand curve does not shrink. It *migrates*. Simple work routes to local and open models. Hard tasks go to the cloud. The frontier is reserved for users whose ROI is provable; Zito names Citadel and Jane Street as gladly paying anything for it. Citrini's rotation map points the same way: value moves to smart routing, observability, edge inference, and good-enough models. Volume keeps growing while the dollars and margins pool somewhere other than where today's valuations assume.

That makes the IPO timing awkward. Anthropic filed its confidential S-1 on June 1, at a reported run-rate of $47 billion (up from $10 billion a year earlier) and a reported valuation near $965 billion. It was the same morning flat-rate pricing ended, and days before its largest customers began rationing its product (Exhibit 8). The labs are monetizing token volume at the moment their customers are learning to minimize it. Underwriters will call that growth. Buyers should ask how much of it survives the routing layer.

EXHIBIT 8

Monetization arrives precisely as customers ration

Anthropic revenue run-rate, $B (reported)

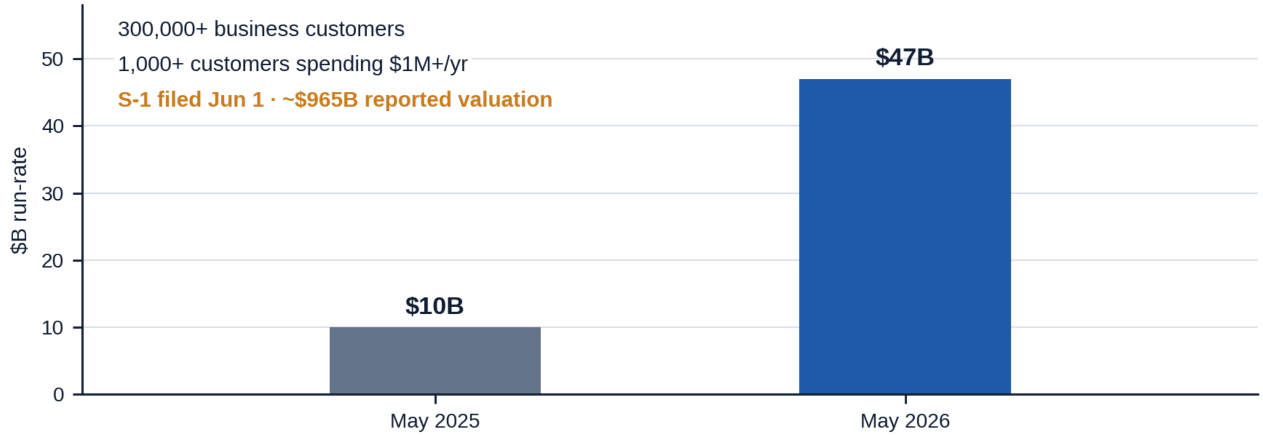

Note: Reported figures, pre-S-1 disclosure; treat as press-reported. Valuation ~$965B per Series H reports.

Source: Fortune / TechCrunch, June 1, 2026 (S-1 reports); The Economist via Citrini (May 2026 run-rate).

WHAT WE'D DO · THE OPERATOR'S PLAYBOOK

1. Re-denominate your metrics. Instrument cost per completed task for your top five AI workflows. Retire "AI usage" KPIs; they buy you workslop and leaderboard-driven token burn (Amazon's lesson), not output.

2. Write a routing policy. Default to the cheapest model that clears your quality bar and escalate to frontier models only on documented ROI. A three-tier ladder (local, mid, frontier) captures most of the deflation Zito is describing.

3. Govern tokens like headcount. Caps, alerts, and dashboards (Uber's $1,500-a-month template), paired with workflow redesign so the caps don't freeze adoption at its least productive configuration.

4. Re-underwrite vendor contracts. Anything signed under flat-rate assumptions needs stress-testing at 2x to 3x usage growth. Mark September 1: the end of GitHub's transition cushion and the first full-price month.

5. Set your labor-to-compute ratio deliberately. The capital swap is happening either way (Exhibit 4). Decide what share of next year's capacity growth is hired versus computed, and own the trade-off explicitly.

WATCHLIST · WHAT WOULD CHANGE OUR MIND

Silicon Data token index: a sustained break lower separates deflation in the spend line from demand destruction. The best real-time tell on the cycle.

September 1: GitHub Copilot's transition pricing ends. The first clean read on metered demand.

Anthropic roadshow: how much revenue is committed spend versus meterable usage that customers are actively learning to reduce.

Q2 earnings (late July): hyperscaler capex guidance against softening token expenditure. The first quarter where the two series can visibly diverge.

Challenger June report (early July): does AI attribution hold above 40%, and does anyone start un-attributing as the narrative cools?

SCORECARD · CONSENSUS VS. THE TAPE

“Token costs only go down.” Wrong at the expenditure level. Usage-weighted spend per token roughly doubled from December to May as agents pushed workloads to frontier reasoning models.

“Enterprise demand is the bubble’s weak link.” Inverted. Demand exploded; value per dollar became the weak link. Rationing began only in Q2 2026.

“95% of pilots fail, so spending will stall.” Spending accelerated into bill shock first. Failure and spend growth coexisted for three full quarters.

NEXT ISSUE · THE ENTRY-LEVEL PUZZLE

Stanford’s payroll data says AI is displacing workers aged 22 to 25 in exposed occupations, a 16% relative employment decline. The NY Fed says 64% of the rise in graduate unemployment comes from remote work, not AI. Yale says the occupational mix isn’t changing at all. Someone is measuring wrong. Next week we show you who, and what it means for your hiring pipeline.

SELECTED SOURCES

Bloomberg interview with John Zito, Morgan Stanley US Financials Conference (Jun 10, 2026) · Citrini Research, “State of the Themes” (Jun 8, 2026) · D. Thompson with D. O’Laughlin, “The AI Boom’s ‘Wait, Is This Worth It?’ Phase” (May 29, 2026) · Axios, “AI sticker shock” (May 28, 2026) · WSJ, “Corporate America Is Starting to Ration AI” (May 2026) · Bloomberg/TechCrunch on Uber caps (Jun 2, 2026) · GitHub changelog (Jun 1, 2026) · UBS Evidence Lab enterprise checks via Bloomberg (Jun 2026) · Silicon Data Token Market Pulse · Stanford AI Index 2025 · Gartner CFO surveys via The CFO (Feb 2026) · Atlanta Fed macroblog (May 6, 2026) · Challenger, Gray & Christmas (Jun 4, 2026) · Apollo 2026 Outlook (T. Slok) · Morgan Stanley via Fortune (May 2026) · P. Kedrosky, “AI Capex Is Eating the Economy” · Fortune/TechCrunch on Anthropic S-1 (Jun 1, 2026) · BIS WP 1179 · G. Gopinath, Odd Lots (May 29, 2026) · HBR, Davenport & Srinivasan (Jan 2026).

Estimates and interpolations are flagged on each exhibit. The $500M single-month bill is single-sourced (Axios) and unconfirmed. Second Order is a client briefing of Razor Research. It is provided for discussion purposes only and is not investment, legal, tax, or accounting advice.