

SECOND ORDER

AI, Capital & Work · A Second Order research briefing

ISSUE 04 · JUNE 30, 2026

The Stack Splits

Over two weeks in June the United States restricted its two most capable public models and gated the next. China released its best model in the same weeks, free, as open weights that run on Chinese chips. The variable that sets the global AI standard has moved from capability, where the US still leads by five points, to diffusion, where US policy is now the obstacle.

FOR DISCUSSION · NOT INVESTMENT ADVICE

Who sets the standard the world builds on

THE THROUGH-LINE

Since Issue 01 the series has read the buildout from inside the US, where the capital and the frontier sit. The events of June force a different question, which is who controls the model the world runs. Washington suspended its two most capable public models and held a third for review over the same two weeks China released a top model as open weights on its own silicon. The capability gap is down to five points among deployable models, and the deployment gap widened in the same weeks.

WHY NOW

The June 12 directive pulled Fable 5 and Mythos 5. GLM-5.2 shipped open on June 13. The White House held GPT-5.6 on June 25. The economics that govern model procurement changed over these two weeks, whether or not the roadmap has caught up.

WHAT YOU WILL LEARN

1. The June split, three US restrictions against one open release (sec. 01)
2. Why the binding constraint is policy, with a five-point gap (sec. 02)
3. How the base layer tilted toward China, 229M against 105M downloads (sec. 03)
4. A full stack on 100,000 Huawei accelerators (sec. 04)
5. Six operator moves, a dated watchlist, and the falsifier

THE ONE LINE · the standard is set at the layer developers build on, and 41% of Hugging Face downloads are now Chinese, a plurality.

The controls meant to slow China are slowing the US

“At a time when access to frontier models is abruptly cut off for non-technical reasons, we are even more convinced.”

z.ai (Zhipu), on releasing GLM-5.2 as open weights, June 13, 2026

51

GLM-5.2 on the Intelligence Index, the top open model, five points under the best deployable US model at 56

2.2x

Chinese open families draw 229M monthly Hugging Face downloads against 105M for US open families (Second Order)

0

Nvidia GPUs in GLM-5's training, which ran on about 100,000 Huawei Ascend 910B accelerators

Three US restrictions against one open release

EXHIBIT 1**Two US model restrictions and one open Chinese release, June 12 to 25**

Frontier-model access events, June 2026



Note: Fable 5 and Mythos 5 access suspended by US export-control directive; GPT-5.6 (Sol, Terra, Luna) held pending review.

Source: Anthropic statement (Jun 12); z.ai/Zipu release (Jun 13); Musk on X (Jun 18); TechCrunch, CNN, Axios (Jun 25).


SO WHAT

- Over two weeks the US restricted three frontier models while China released one, GLM-5.2, as open weights.
- The June 12 directive targeted foreign-national access to Fable 5 and Mythos 5, so with no way to segment accounts Anthropic shut both for everyone.
- The June 25 White House hold put GPT-5.6 into a review of up to 30 days before release.

The binding constraint moved to policy

EXHIBIT 2**The capability gap is small, and the #1 model is the one the US pulled**

Artificial Analysis Intelligence Index, leading models, as of Jun 27 2026



Note: Fable 5 (60) leads but is suspended in the US; GLM-5.2 is the top open-weight model, on a Huawei stack.

Source: Artificial Analysis Intelligence Index (artificialanalysis.ai), retrieved Jun 29 2026.


SO WHAT

- The best deployable US model, Opus 4.8 at 56, leads GLM-5.2 at 51 by five points on the Intelligence Index.
- Fable 5 at 60 sat at the top of that table, and the June 12 directive pulled it.
- A five-point edge among deployable models is thin enough for an installed base to erode, and that is under way.

Open weights compound, and adoption has tilted

EXHIBIT 3**The open-weight base has flipped to China**

Monthly Hugging Face downloads, leading open LLM families, top models per family (millions)



Note: Sum of each family's top models by 30-day downloads; counts overstate real usage (automated and CI pulls).

Source: Second Order calculations from the Hugging Face Hub API (huggingface.co/api/models), retrieved Jun 29 2026.


SO WHAT

- Chinese open families draw about 229M Hugging Face downloads a month against about 105M for US open families (Second Order, June 29).
- Qwen alone draws about 180M, more than every US open family combined, and anchors 113,000 fine-tuned derivatives, over 200,000 counting every model that tags it, above Google and Meta together.
- Open weights are about a third of OpenRouter tokens, and Chinese models are growing fastest.

The stack runs on 100,000 Huawei accelerators

EXHIBIT 4**The Huawei stack runs about a third slower than Nvidia**

Reported GLM-5.2 inference throughput vs Nvidia-backed peers (tokens/sec)



Note: Inference rates are reported ranges (whiskers), not benchmarked here; the training-chip claim is vendor-originated.

Source: Training-chip claim: digitaltoday.co.kr, letsdatascience.com, techtimes.com (multi-outlet). Inference rates: reported, single-source.

SO WHAT

- GLM-5 trained on about 100,000 Huawei Ascend 910B accelerators on the MindSpore framework, with no Nvidia hardware (vendor-reported).
- The cost is a throughput penalty near a third, about 17-19 tokens/sec against 25-30 for Nvidia-backed peers (single source, shown as a range).
- Open weights, Chinese silicon, and a platform where Chinese models lead form a stack outside US jurisdiction at every layer.

Standards are set by what developers build on

The first-order reading of the June events favors Washington, since keeping the best models at home denies them to rivals and protects the US lead.

Standards are won at the base layer, where availability matters more than a peak score, and they compound. Each week a model is the one developers can hold and fine-tune, the evaluation harness, the tooling, and the next layer of work settle onto it, and every team that builds on GLM-5.2 raises its own cost of leaving.

A developer who has watched a frontier model go dark by directive prices that risk now, and prices it toward weights that cannot be recalled, which today means the Chinese families. The US lead in raw capability is real, and each month of restricted access converts a little of it into installed base for the open Chinese models.

Six moves for operators

1. Treat access as a supply line

A model reachable only through a vendor a directive can shut is a single point of failure. Name a fallback model and the trigger that moves you to it.

2. Test an open fallback now

Keep a self-hostable open model in parallel to production so the switching cost is a measured quantity before you need it under pressure.

3. Split weights from jurisdiction

Run open weights on infrastructure you control, and keep regulated data off any China-hosted endpoint.

4. Build model-agnostic

Standardize on interfaces that swap the underlying model within a day, because the leaderboard and the policy regime will keep moving.

5. Map your silicon

Know your exposure to Nvidia against the alternatives, and to the memory bottleneck behind all of it. Next week takes up that capex.

6. Assume open abroad

Customers outside the US will build on open, increasingly Chinese, weights. Price and integrate for that as the default.

WATCHLIST

What would change our mind

GPT-5.6 review · A broad release within weeks marks the June 25 hold as a delay, and a long hold marks it as a standing license over frontier deployment.

Fable 5 restored July 1 · Anthropic returns Fable 5 worldwide today, which argues the June 12 suspension worked as a narrow fix and returned the model within three weeks.

Q1 2027 horizon · Musk and z.ai both name the first quarter of 2027 for an open model at Fable's level. Watch the Intelligence Index gap close or hold.

Open-weight share · Watch whether GPT-OSS, Gemma, and OLMo regain download and derivative share on Hugging Face and OpenRouter.

The falsifier · A fast, durable US carve-out that ships frontier models worldwide, plus a plateau in Chinese open adoption, would break the thesis.

NEXT ISSUE · THE CAPEX ECONOMY. · Growth is increasingly the buildout itself, and a rising share of it is memory. Gavin Baker puts high-bandwidth memory at 30 to 40% of hyperscaler capex by 2027, supplied by a three-firm oligopoly, two of them Korean. We separate the capex figure markets quote from the capacity it buys.