

SECOND ORDER

AI, Capital & Work · A Second Order research briefing

DESK SIDEBAR · JUNE 15, 2026

The 7% Problem

Companies spend 93 cents of every AI dollar on the tools and 7 on the people who make the tools pay off. That 7 cents is where the missing AI returns are hiding.

The unspent 7% is the missing AI return

THE THROUGH-LINE

The Star Tribune piece reports a 93/7 budget split and, separately, that few firms can show a return on AI. This sidebar reads them as one fact, since the 7% that firms decline to spend on enablement is the complementary investment that turns tools into measured output.

WHY NOW

AI use is near-universal while enablement spend lags. Hiring plans and budgets get set this quarter, and the firms pulling ahead are the ones funding the 7% while their peers buy bigger models.

WHAT YOU'LL LEARN

1. The 7% is the ROI gap, via the J-curve (\$01)
2. The trainer wall, one per 312 workers (\$02)
3. Shadow AI is the release valve (\$03)
4. How the winners integrate training (\$04)
5. The skill premium shifts, plus six moves

ONE LINE TO CARRY the firms that book a return on AI are the ones funding the 7%.

The training gap is the missing-returns story

Companies spend about 93% of AI budgets on technology and 7% on people and process. That 7% is the complementary investment the productivity literature says must come first, and skipping it is why the returns have not shown up.

The firms that book a return on AI are the ones funding that 7%.

95%

of enterprise generative-AI pilots show no measurable profit impact (MIT)

1 : 312

dedicated trainers to workers. You cannot classroom your way to AI fluency (BLS)

44% / 41%


use AI against policy vs. work somewhere that has a policy at all (KPMG)

The unspent 7% is the AI return itself

EXHIBIT 1

93 cents of every AI dollar buys tools and 7 buys the people who make them pay

Share of corporate AI budget (left) and how far AI initiatives reach (right)



Note: Second Order compilation. 93/7 is single-source Deloitte, first framed around finance teams; treat as directional. Funnel shares are of organizations.

Source: Deloitte (93/7 split); MIT NANDA, State of AI in Business 2025 (funnel); McKinsey State of AI 2025. Data: data/budget_vs_payoff.csv.

SO WHAT



- 93% of AI budgets go to tools, 7% to people and process (Deloitte; directional).
- 80% explore, 20% pilot, and 5% reach measurable P&L impact (MIT).
- The 7% is the complementary investment the J-curve says comes first.

You cannot classroom your way to AI fluency

EXHIBIT 2

About one dedicated trainer serves every 312 U.S. workers

U.S. dedicated workplace trainers against the workforce they would reach



Note: 2019 is OEWS, 2024 the broader Employment Projections base year; not strictly comparable. Ratio = 2024 trainers over May 2025 jobs. AI-user count is an estimate.

Source: BLS OEWS (May 2019, May 2025); OOH 2024-34 projections (SOC 13-1151, 11-3131). Second Order calculations. Data: data/trainer_capacity.csv.


SO WHAT

- ~499,000 dedicated trainers for 155.5M jobs, one per 312 workers (BLS).
- Growing fast (specialists +11%) and still adds only ~51,500 by 2034.
- ~249 AI users per trainer. The 7% cannot run through a central L&D team.

Skip the training and the AI goes underground

EXHIBIT 3**87% of U.S. workers use AI weekly and 41% work under a written policy**

U.S. employees and the gap between AI use and AI governance, 2025



Note: KPMG figures are the primary anchors; 80% unapproved-tools and 38% sensitive-data readings are secondary. Bars show share of employees except the bottom bar (share of organizations).

Source: KPMG U.S., 2025 (Shadow AI is already here; American Worker survey). Second Order compilation. Data: data/shadow_ai.csv.

SO WHAT

- 87% use AI weekly; 44% against policy; 58% never check the output (KPMG).
- Only 41% work somewhere with a generative-AI policy at all.
- The training-budget gap converts into a data-governance gap.

The firms with a return fund the 7%, in the line

EXHIBIT 4
The firms that book a return on AI fund the 7%

Two operating patterns behind the same tools

	TOOL-FIRSTno measured return	CAPABILITY-FIRSTthe 7% gets funded
Sequence	Buy the tool, roll it out, hope people figure it out	Train first, then roll out (Affinity Plus: Copilot after org-wide training)
Direction	Top-down; end users not consulted	Co-designed with the people who use it (Mayo: nursing AI contest, staff use cases)
Training scope	Generic access; a license for everyone	Role-specific: how to use, when NOT to, how to verify the output
Guardrails	None; "shadow AI" fills the vacuum	Explicit policy plus peer forums and office hours (Thrivent: >90% monthly active)
How it is measured	Seats deployed; tools purchased	Adoption, hours saved, and P&L (advisers save ~4 hrs/week)

Source: Second Order synthesis of reported firm practice and McKinsey/BCG research; framework, not a computed result.

SO WHAT

- Train before rollout; teach when NOT to use it and how to verify (Affinity Plus).
- Enablement runs peer-led and embedded through the line (Thrivent, >90% active).
- Co-designed with users; measured on adoption, hours, and P&L (Mayo).

Six moves to fund the 7% this quarter

1. Re-cut the budget

Move from 93/7 toward 80/20. Once the tool is good enough, the next dollar of enablement beats the next model upgrade.

2. Train before you deploy

Teach the negative space, when to avoid the tool and how to verify it. 58% of users skip that check.

3. Push enablement to the line

At one trainer per 312 people, a central academy cannot be staffed, so scale with peer champions and office hours.

4. Policy before spread

Coverage at 41% is the vacuum shadow AI fills. Write the policy before the tool spreads, not after.

5. Measure adoption, then P&L

Seats deployed measures procurement. Monthly-active near 90% and ~4 hours saved per person measure the payoff.

6. Hire for judgment

Screen for the instinct to verify and the ability to coach a tool through a team, and weight it above raw technical fluency.

WATCHLIST

What would change our mind

EBIT among tool-heavy spenders. If enterprise earnings impact climbs without a shift in enablement spend, the thesis weakens.

Self-adapting agents. MIT blames a learning gap. If agents start adapting to workflows, the 7% matters less.

The falsifier. tool-heavy, training-light firms booking more than 5% of EBIT from AI. We have not found it.

THE FIX IS STRUCTURAL. At one trainer per 312 workers, AI fluency cannot come through L&D headcount, so the teaching moves into the line and hiring shifts toward judgment.